Doc 4: CDQN Concept Proposal Paper - Version V2.0.0 (Enhanced Concept Proposal)

Metadata:

- **Version:** V2.0.0
- **Date:** 2025-07-13T04:20:00Z
- **Agent:** Gemini: Google (2025-07-09)
- **Lead Author:** Christophe Duy Quang Nguyen
- **Human Contributors:** N/A
- **License:** BaDaaS License - The Agile Commercial Open-Core License (Doc 2)
- **Changelog:** Updated to V2.0.0. Comprehensive revision to integrate advanced concepts from recent research in memory management, multi-agent systems, LLM training, and communication protocols, enhancing the CDQN project's foundational architecture and addressing specific design questions for a more robust and self-evolving AI ecosystem.
- **Summary:** This document outlines the foundational concepts and architecture for the "CDQN" project, detailing its components "cdqnDB," "cdqnProtocol," and "cdqnLang," and presenting its vision for ubiquitous, affordable AI through the intelligent orchestration and management of self-evolving nano language models with context data.
- **Sections Affected:** All sections.
- **Contact:** cdqn5249@gmail.com

1. Introduction: The Need for Ubiquitous, Affordable Intelligence

The current paradigm of Artificial Intelligence (AI) is heavily skewed towards large, centralized cloud-based Large Language Models (LLMs), which demand immense computational resources and impose significant latency, privacy, and cost barriers. This model fundamentally limits AI's ubiquity, especially in remote, resource-scarce, or privacy-critical environments.

The **CDQN (Context Datas Quorum Nodes) project** is proposed as a transformative paradigm shift, leveraging the collective intelligence of numerous, highly specialized, and extremely efficient "nano agents" to achieve emergent capabilities comparable to, or even surpassing, those of large models in specific contexts. This vision is deeply informed by natural systems, where complex behaviors arise from the simple interactions of many individual components.

The **cdqNetwork** represents the overall swarm structure of interacting nano-agents and nodes within the CDQN framework. It aims to provide a standardized, decentralized framework for building and deploying specialized nano-agents that can operate on low-power, embedded systems.

2. The CDQN Architecture: Swarm Intelligence Meets AI

CDQN's core architecture is a **decentralized multi-agent system** built upon three foundational pillars: **cdqnDB**, **cdqnProtocol**, and **cdqnLang**.

This design ensures that detection nodes operate collaboratively, share minimal metadata, and make localized decisions, thereby avoiding bottlenecks related to centralization.

2.1. cdqnDB: The Evolving Collective Database with FSSF Layers

cdqnDB is envisioned as a new type of distributed, evolving database that serves as the collective memory and storage for the entire cdqNetwork. It enables agents to efficiently store, query, and retrieve "Context Data".
This database is designed to manage the lifecycle of data, including storage, retrieval, compression, pruning, and potential "fusion" of redundant or complementary knowledge, mimicking synaptic plasticity.

The core unit of storage in cdqnDB is the **Context Data Unit**, which is a reproduction and extension of the **MemCube** concept from the MemOS framework. A Context Data Unit encapsulates not just content (e.g., plaintext, sensor data, model weights, activation states, structured graphs, or prompt templates) but also rich metadata.
This metadata includes critical attributes such as **provenance**, **confidence scores**, **temporal tags**, **semantic context**, **access permissions**, **versioning**, **last usage time**, **source model**, **usage frequency**, **priority**, **expiration dates**, **embedding fingerprints**, and **storage mode** (e.g., compressed). This detailed metadata enables structured and queryable knowledge graph formation within the swarm.

To enhance knowledge management and address concerns like hallucination, cdqnDB introduces **FSSF (Factual, Semi-factual, Semi-fiction, and Fiction) sub-database layers**. These layers are implemented as a semantic classification within the metadata of Context Data Units, leveraging the flexible metadata.tags or payload.format attributes as per the MemCube design. This classification allows cdqnDB to manage different types of knowledge based on their veracity and origin, enabling fine-grained control and retrieval based on trustworthiness. The system can prioritize factual data for critical reasoning tasks while allowing for creative or hypothetical contexts when appropriate. This multi-level classification is supported by MemOS's MemGovernance features, including access control, versioning, and provenance auditing, ensuring secure, trustworthy, and compliant intelligent agent systems.

The cdqnDB's evolving schema, which mimics synaptic plasticity, will be realized through MemOS's principles of **plasticity** and **evolvability**. This allows for dynamic restructuring, migration, and seamless transitions between different memory types—including parameter memory (knowledge embedded in model weights), activation memory (contextual inference state, such as KV caches), and plaintext memory (external structured knowledge). This adaptive approach ensures that the database continuously refines its organization based on the learning and evolving needs of the agents.

For performance optimization, cdqnDB integrates concepts from KV cache management systems like LMStack and Mooncake. This includes **decoupling KV caches** from serving engines to provide better performance and manageability, and implementing **dynamic compression, composition, and modification of KV caches** through a Knowledge Delivery Network (KDN). This enhances global cache capacity and inter-node transfer bandwidth for nano agents.

Furthermore, cdqnDB's querying mechanisms will be informed by graph-based RAG approaches such as NodeRAG and GraphRAG-Causal, enabling **hybrid search capabilities** that combine semantic embeddings with structural cues for highly relevant and contextually aware retrieval.

2.2. cdqnProtocol: The Multi-layered Agentic Communication Layer

cdqnProtocol defines the communication paradigms for the cdqNetwork, drawing inspiration from nature to facilitate efficient and intelligent inter-agent communication. It is conceived as a multi-layered protocol, each layer addressing different aspects of agent interaction and data flow, enhancing the original bio-inspired concepts.

1. **Orchestration & Coordination Layer:** This highest layer is responsible for high-level task distribution, dynamic agent selection, and overall workflow management for complex multi-agent tasks. It will implement agent orchestration mechanisms to automatically optimize task scheduling and prompt generation, thereby reducing communication overhead and saving computing resources for the involved nano agents.
2. **Semantic Communication & Trust Layer:** This layer focuses on enabling semantically rich and context-aware message exchange among agents, incorporating ethical annotations as seen in the LOKA protocol. It will manage unified content provenance and trust frameworks to ensure the reliability and traceability of knowledge shared between agents, mitigating risks associated with compromised or outdated information.
3. **Bio-Inspired Primitives Layer:** This layer retains the original bio-inspired elements proposed in CDQN, such as pheromone_signal and SwarmSearch. These primitives will facilitate emergent collective behaviors, utilizing cdqnDB for the dynamic creation and decay of "digital pheromone trails" that guide agent interactions and information propagation within the network.
4. **Low-Level Data Transfer Layer:** Designed for efficient and high-throughput data movement between nodes, this layer draws inspiration from technologies like RDMA networks used in Mooncake and the principles of disaggregated memory systems seen in Beehive. It ensures rapid and reliable data exchange, critical for low-latency operations in distributed edge environments.

The cdqnProtocol will proactively integrate **security and safety countermeasures** throughout its layers, drawing from surveys on agent communication security. This includes implementing data minimization, multi-layered input validation, output consensus filtering, sandboxing of tool execution, and continuous monitoring to defend against privacy leakage, OverThink attacks (excessive token consumption), and tool-related vulnerabilities.

2.3. cdqnLang: The New Object-Context-Oriented Programming Language

To effectively build and manage the CDQN framework, **cdqnLang** is proposed as a novel programming language. It blends the performance and safety characteristics of Rust and Go with a novel, context-aware syntax that naturally supports agentic behavior and swarm interactions.
The defining innovation of cdqnLang is its **object-context-oriented programming paradigm**. This means that cdqnLang natively treats **Context Data Units** (MemCubes) as first-class objects within its syntax. Operations for interacting with cdqnDB, such as cdqnDB.publish_context_data(type, content, metadata) and cdqnDB.query(filter_criteria), are not merely API calls but are core language features, abstracting away complex distributed data management and making interaction with the collective memory intuitive and efficient.

Inspired by CodeAgents, cdqnLang will facilitate **codified reasoning** by enabling nano agents to express their multi-agent interactions, planning, execution, and feedback in modular pseudocode, enriched with control structures (e.g., loops, conditionals), boolean logic, and typed variables. This approach significantly enhances the interpretability, reduces ambiguity, and improves the traceability and token efficiency of multi-step reasoning processes.

cdqnLang will feature **Actor-Model First** support, providing strict type-safe, isolated, and message-passing agents, similar to Erlang. This paradigm is crucial for concurrency and distribution within the cdqNetwork.
Furthermore, cdqnLang will incorporate an **Adaptive Compilation** or "living compiler". This compiler will dynamically optimize nano agent code on the fly based on the deployment environment and observed agent behavior, using feedback loops. This is akin to ALTA's compiler that maps programs to Transformer weights and AuDaLa's ability to provide practical compilation for fast parallel execution on various hardware.

3. Training and Self-Evolution of Nano Agents

The **nano agents** within the cdqNetwork are designed to be highly specialized and extremely efficient, capable of running on microcontrollers, edge AI chips, and low-power CPUs. Their training and self-evolution will be driven by a synergistic combination of architectural innovation and advanced learning paradigms.

Mixture-of-Experts (MoE) Architecture:
A foundational aspect for the scalability and specialization of nano agents will be the widespread adoption of **Mixture-of-Experts (MoE)** architectures. MoEs allow for a dramatic expansion of model capacity without a proportional increase in per-token compute, enabling nano agents to become highly specialized while remaining computationally efficient. This approach routes inputs to specific "experts" (sub-models), promoting their specialization. Examples like DeepSeekMoE, Mistral 8x7B, Kimi-VL, and Jamba demonstrate the effectiveness of MoE in achieving high performance with efficient parameter activation [cite: 1, 6, 43, 53-55].

Reinforcement Learning (RL) for Self-Evolution:
Reinforcement Learning will be the primary mechanism enabling nano agents to self-evolve and continually improve at their tasks. Inspired by Self-Adapting LLMs (SEAL), nano agents will be trained to generate their own finetuning data and update directives. The downstream performance of the updated model will serve as the reward signal for the RL loop, allowing agents to persist weight updates and adapt to new knowledge. SAGE-nano's concept of inverse reasoning further enables nano agents to introspect on their own reasoning processes, providing metacognitive self-reflection and enhancing interpretability alongside reasoning performance. This allows nano agents to understand not just *what* they decided, but *why*.

Iterative Refinement and Feedback Loops:
Continuous self-correction and improvement will be integral to the nano agents' self-evolution. This includes:

- **Concept-Driven Specialization:** Nano agents will be trained with deep conceptual knowledge relevant to their domains (similar to UniTranslator's "concept agent" fine-tuned on programming language data). This ensures they develop a profound understanding of

their tasks, moving beyond mere surface-level syntax.
- **Compiler and NLI Feedback:** Drawing from UniTranslator, nano agents will leverage Natural Language Inference (NLI) grounding and compiler feedback mechanisms for real-time code validation and error mitigation, refining their output iteratively.
- **Failure-Driven Adjustments:** Inspired by SwarmAgentic, nano agents will incorporate failure-driven adjustments, personal best guidance (self-learning), and global best guidance (swarm-based learning) into their optimization processes, iteratively refining their functionalities and collaboration strategies. CodeAgents' use of in-plan assert checks and global feedback loops will enable robust error handling and replanning on critical failures.
- **Addressing Multi-Agent System (MAS) Failures:** The self-evolution mechanism will explicitly incorporate strategies to mitigate MAS failures. This includes promoting clear role specifications, fostering modular agent design, and utilizing cross-verification mechanisms (e.g., multiple LLM calls with majority voting) to enhance reliability and resilience, as highlighted by the MAST taxonomy.

Distributed Edge AI Deployment:
Nano agents are fundamentally designed for operation in resource-constrained, heterogeneous edge environments. The CDQN framework will leverage an adaptive split inference orchestration framework to dynamically manage LFM partitions across edge nodes. [cite_start]This approach enables **real-time, Quality-of-Service (QoS)-aware management** of inference workloads and enhances **privacy preservation** by localizing sensitive computations directly on the edge, reducing data movement to centralized clouds [cite: 6, 67-70].

4. Conclusion: Towards Symbiotic AI with CDQN

The CDQN framework, powered by its evolving **cdqnDB**, sophisticated **cdqnProtocol**, and native **cdqnLang**, represents a bold leap towards a future where sophisticated AI is no longer confined to data centers but becomes an integral, ubiquitous, and affordable part of our everyday lives and critical infrastructure. By learning from nature's unparalleled efficiency in distributed intelligence and self-organization, CDQN unlocks a new era of AI that is more robust, private, adaptable, and truly symbiotic with its environment.

The enhanced design addresses key questions:
- **CDQN (Context Datas Quorum Nodes) project:** A paradigm shift for ubiquitous, affordable AI through collective intelligence of specialized nano agents, powered by a decentralized multi-agent system.
- **cdqNetwork:** The overall swarm structure of interacting nano agents and nodes within the CDQN framework, enabling standardized, on-device AI deployment.
- **Training and self-evolution of nano LMs:** Achieved through **Mixture-of-Experts (MoE)** architectures for efficiency and specialization, primarily driven by **Reinforcement Learning (RL)** that allows agents to generate their own finetuning data and self-reflect using inverse reasoning. Continuous **iterative refinement** with concept-driven specialization, compiler/NLI feedback, and failure-driven adjustments ensures ongoing improvement and resilience.
- **Detailed specifications of cdqnDB, cdqnProtocol, and cdqnLang:**
    - **cdqnDB:** An evolving, distributed, collective memory, with **Context Data Units** (akin to MemCubes) encapsulating rich content and metadata (including FSSF

layers for semantic classification). It supports dynamic data management, knowledge graph formation, and leverages KV cache optimization for performance.
    - **cdqnProtocol:** A multi-layered communication system including **Orchestration & Coordination**, **Semantic Communication & Trust**, **Bio-Inspired Primitives**, and **Low-Level Data Transfer layers**. It integrates robust security countermeasures against various agent-related risks.
    - **cdqnLang:** An **object-context-oriented programming language** that natively treats Context Data Units as first-class objects, enabling codified reasoning with structured pseudocode, supporting an Actor-Model First paradigm, and featuring an Adaptive Compilation or "living compiler" for on-the-fly optimization.

The modular architecture facilitates seamless integration with existing security frameworks and supports horizontal scaling in distributed data centers. By treating memory as a system-level resource and fostering a self-evolving ecosystem, CDQN sets the foundation for truly intelligent, autonomous, and accountable AI.

5. Updated Glossary

- **CDQN (Context Datas Quorum Nodes):** The proposed project aiming for ubiquitous, affordable AI via collective intelligence of specialized nano agents in a decentralized multi-agent system.
- **cdqNetwork:** The overall swarm structure of interacting nano-agents and nodes within the CDQN framework.
- **cdqnDB:** The proposed evolving, new type of distributed database and collective memory for the CDQN network, designed to store and manage "Context Data".
- **cdqnLang:** The proposed new object-context-oriented programming language specifically designed for developing and managing agents within the CDQN framework.
- **cdqnProtocol:** The proposed set of communication concepts and protocols (including bio-inspired and multi-layered ones) for inter-agent communication within the CDQN framework.
- **Context Data Unit:** The atomic unit of storage in cdqnDB, encapsulating content (e.g., text, sensor data) along with rich metadata. It is analogous to a MemCube and enables structured and queryable knowledge within the CDQN system.
- **FSSF (Factual, Semi-factual, Semi-fiction, Fiction):** Proposed semantic classification layers within the metadata of Context Data Units in cdqnDB, designed to manage knowledge based on veracity and origin, enhancing knowledge management and reducing hallucination risks.
- **Inverse Reasoning:** A novel paradigm enabling LLMs to decompose and explain their own reasoning chains post-hoc, providing insight into why specific reasoning choices were made.
- **MemCube:** A basic unit of memory in the MemOS framework, encapsulating memory content and metadata such as provenance and versioning, serving as inspiration for cdqnDB's Context Data Units.
- **Mixture-of-Experts (MoE):** A neural network architecture that dramatically expands model capacity and achieves remarkable performance without increasing per-token compute by routing inputs to specialized expert sub-networks.
- **Nano Agent:** Highly specialized and efficient AI agents designed to run on microcontrollers, edge AI chips, and low-power CPUs as part of the cdqNetwork.

- **Object-Context-Oriented Programming:** A proposed programming paradigm for cdqnLang where Context Data Units (MemCubes) are treated as native, first-class objects, allowing direct and intuitive interaction with a distributed knowledge base.
- **Reinforcement Learning (RL):** A machine learning paradigm where an agent learns to make decisions by performing actions in an environment to maximize a cumulative reward, used for the self-evolution of nano agents.

6. References (Excluding CDQN Documents)

- [73] Nafea Alhamadi. Distributed Denial of Service (DDoS) Flooding Attack Mitigation Using Ensemble Deep Learning and Multi-Agent Systems for Cloud Computing Environments. *Journal of Soft Computing and Data Mining*, 6(1):406–422, 2025.
- [43] Abubakr O Al-Abbasi, Vaneet Aggarwal, and Moo-Ryong Ra. Multi-tier caching analysis in cdn-based over-the-top video streaming systems. *IEEE/ACM Transactions on Networking*, 27(2):835–847, 2019.
- [74] Google Cloud. Vertex AI Agent Engine Memory Bank overview. 2024. URL: https://cloud.google.com/vertex-ai/docs/agent-engine/memory-bank.
- [75] Yihua Cheng. *Optimizing Distributed Large Language Model Serving by Decoupling KV Caches*. Dissertation, Carnegie Mellon University, 2024.
- [47] Peter Shaw, James Cohan, Jacob Eisenstein, Kenton Lee, Jonathan Berant, and Kristina Toutanova. ALTA: Compiler-Based Analysis of Transformers. arXiv preprint arXiv:2410.18077, 2024.
- [61] Rabimba Karanjai, Sam Blackshear, Lei Xu, and Weidong Shi. COLLABORATION IS ALL YOU NEED: LLM ASSISTED SAFE CODE TRANSLATION. arXiv preprint arXiv:2503.11237, 2025.
- [65] Anthropic, Google, and NVIDIA. Why Do Multi-Agent LLM Systems Fail? An Empirically Grounded Taxonomy and Structural Interventions. arXiv preprint arXiv:2503.13657, 2025.
- [67] Fernando Koch, Aladin Djuhera, and Alecio Binotto. Intelligent Orchestration of Distributed Large Foundation Model Inference at the Edge. arXiv preprint arXiv:2504.03668, 2025.
- [53] Kimi Team. KIMI-VL TECHNICAL REPORT. arXiv preprint arXiv:2504.07491, 2025.
- [29] Babar M, Ali Abbas, and Insaf Ullah. GraphRAG-Causal introduces an innovative framework that combines graph-based retrieval with large language models to enhance causal reasoning in news analysis. arXiv preprint arXiv:2506.11600, 2025.
- [51] Houyi Li, Ka Man Lo, Ziqi Wang, Zili Wang, Wenzhen Zheng, Shuigeng Zhou, Xiangyu Zhang, and Daxin Jiang. Can Mixture-of-Experts Surpass Dense LLMs Under Strictly Equal Resources? arXiv preprint arXiv:2506.12119, 2025.
- [76] MiniMax. MiniMax-M1: Scaling Test-Time Compute Efficiently with Lightning Attention. arXiv preprint arXiv:2506.13585, 2025.
- [64] Haoran Xu, Yixiao Ge, and Tianjun Bu. SwarmAgentic: A Particle Swarm Optimization Approach for Automated Agentic System Generation. arXiv preprint arXiv:2506.15672, 2025.
- [77] Tom T.P. Franken, Thomas Neele, and Jan Friso Groote. The Autonomous Data Language – Concepts, Design and Formal Verification. arXiv preprint arXiv:2506.19457, 2025.
- [38] Dezhang Kong, Shi Lin, Zhenhua Xu, Zhebo Wang, Minghao Li, Yufeng Li, Yilun Zhang, Hujin Peng, Zeyang Sha, Yuyuan Li, Changting Lin, Xun Wang, Xuan Liu, Ningyu

Zhang, Chaochao Chen, Muhammad Khurram Khan, and Meng Han. A Survey of LLM-Driven AI Agent Communication: Protocols, Security Risks, and Defense Countermeasures. arXiv preprint arXiv:2506.19676, 2025.

- [18] Haojun Wu, Tongya Lu, Wen-Chih Lo, Yujie Li, Zhen Guo, Jun Yang, and Min Yen Kan. ReG: Reorganizing Retrieval for Large Language Models. arXiv preprint arXiv:2506.22518, 2025.
- [78] Alan Dao and Dinh Bach Vu. Jan-nano Technical Report. arXiv preprint arXiv:2506.22760, 2025.
- [59] Basab Jha, Firoj Paudel, Ujjwal Puri, Zhang Yuting, Choi Donghyuk, and Wang Junhao. Thinking About Thinking: SAGE-nano's Inverse Reasoning for Self-Aware Language Models. arXiv preprint arXiv:2507.00092, 2025.
- [14] Yuwei Hu, Runyuan Zhao, Yanqi Zhang, Haibo Chen, Jiao Zhang, Yanggang Li, Chenxi Wang, Yanwei Hu, Kaiya He, Weize Zhong, Yulin Zheng, Junzhe Ma, Zhaozhuo Xu, Yushi Li, Jiajin Zhang, Yuanze Zhang, Zhenzhong Lan, and John C.S. Lui. MemOS: A Memory Operating System for Large Language Models. arXiv preprint arXiv:2507.03724, 2025.
- [44] Bruce Yang, Xinfeng He, Huan Gao, Yifan Cao, Xiaofan Li, and David Hsu. CodeAgents: A Token-Efficient Framework for Codified Multi-Agent Reasoning in LLMs. arXiv preprint arXiv:2507.03254, 2025.
- [79] DTCP. Agent Infrastructure Preview. 2025. URL: https://www.dtcp.capital/fileadmin/DTCP/Bilder/News/202507_DTCP_Agent_Infrastructure_Preview_v1.pdf.
- [80] Christophe Duy Quang Nguyen. Fundamentals of Secure AI Systems with Personal Data. EDPB, 2025. URL: https://www.edpb.europa.eu/system/files/2025-06/spe-training-on-ai-and-data-protection-technical_en.pdf.
- [56] Adam Zweiger, Jyothish Pari, Han Guo, Ekin Akyürek, Yoon Kim, and Pulkit Agrawal. Self-Adapting Language Models. arXiv preprint arXiv:2506.10943v1, 2025.
- [81] Ruoyu Qin, Zheming Li, Weiran He, Jialei Cui, Feng Ren, Mingxing Zhang, Yongwei Wu, Weimin Zheng, and Xinran Xu. Mooncake: Trading More Storage for Less Computation — A KVCache-centric Architecture for Serving LLM Chatbot. *Proceedings of the 23rd USENIX Conference on File and Storage Technologies*, 2025.
- [37] Quanxi Li, Hong Huang, Ying Liu, Yanwen Xia, Jie Zhang, Mosong Zhou, Xiaobing Feng, Huimin Cui, Quan Chen, Yizhou Shan, and Chenxi Wang. Beehive: Accelerating Disaggregated Memory with Compiler-Managed Asynchrony in Multithreaded Programs. *USENIX Conference on File and Storage Technologies*, 2025.
- [52] Ashu Kalra, Ankit Singh Rawat, Harsh Rawat, Saiful Haq, Thomas T Joshi, Harsh Trivedi, Sri Vardhamanan, Yuandong Tian, and Manzil Zaheer. A Comprehensive Survey of Mixture of Experts. arXiv preprint arXiv:2408.0583, 2024.